

SummaryPresentation

September 20, 2016

```
In [2]: import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import pandas as pd
# import sys
# sys.path.append("BayesianOptimization/")
from bayes_opt import BayesianOptimization
from matplotlib import gridspec
sns.set_context('talk')
sns.set_style('whitegrid')
%matplotlib inline

def posterior(bo, xmin=-2, xmax=10):
    xmin, xmax = -2, 10
    bo.gp.fit(bo.X, bo.Y)
    mu, sigma2 = bo.gp.predict(np.linspace(xmin, xmax, 1000).reshape(-1, 1))
    return mu, np.sqrt(sigma2)

def plot_gp(bo, x, y):

    fig = plt.figure(figsize=(16,8))
    fig.suptitle('Gaussian Process and Utility Function After {} Steps'.format(bo.iterations))

    gs = gridspec.GridSpec(2, 1, height_ratios=[3, 1])
    axis = plt.subplot(gs[0])
    acq = plt.subplot(gs[1])

    mu, sigma = posterior(bo)
    axis.plot(x, y, linewidth=3, label='Target')
    axis.plot(bo.X.flatten(), bo.Y, 'D', markersize=8, label=u'Observations')
    axis.plot(x, mu, '--', color='k', label='Prediction')

    axis.fill(np.concatenate([x, x[::-1]]),
              np.concatenate([mu - 1.9600 * sigma, (mu + 1.9600 * sigma)[::-1]]),
              alpha=.6, fc='c', ec='None', label='95% confidence interval')

    axis.set_xlim((-2, 10))
```

```

axis.set_ylim((0., 1.5))
axis.set_ylabel('f(x)', fontdict={'size':20})
axis.set_xlabel('x', fontdict={'size':20})

utility = bo.util.utility(x.reshape((-1, 1)), bo.gp, 0)
acq.plot(x, utility, label='Utility Function', color='purple')
acq.plot(x[np.argmax(utility)], np.max(utility), 'd', markersize=15,
        label=u'Next Best Guess', markerfacecolor='gold',
        markeredgecolor='k', markeredgewidth=1, clip_on=False)
acq.set_xlim((-2, 10))
acq.set_ylim((0, np.max(utility) + 0.5))
acq.set_ylabel('Utility', fontdict={'size':20})
acq.set_xlabel('x', fontdict={'size':20})

axis.legend(loc=2, bbox_to_anchor=(1.01, 1), borderaxespad=0.)
acq.legend(loc=2, bbox_to_anchor=(1.01, 1), borderaxespad=0.)

```

```

In [29]: from traitlets.config.manager import BaseJSONConfigManager
path = "E:/ProgramFiles_w/Engineering/Anaconda/etc/jupyter/nbconfig"
cm = BaseJSONConfigManager(config_dir=path)
# # cm.update('livereveal', {
# #     'width': 960,
# #     'height': 700
# # })
#
#     'theme': 'sky',
#     'transition': 'zoom',
cm.update('livereveal', {'start_slideshow_at': 'selected'})

# from IPython.core.display import display, HTML
# display(HTML("<style>.container { width:80% !important; }</style>"))

```

```

Out[29]: {'start_slideshow_at': 'selected'}

```

1 Thurston Sexton

Sandia NL Interview
9/27/2016

- Optimization
- Numerical Modeling
- Machine Learning

1.1 Learning Human Search Strategies

1.1.1 From a Crowdsourcing Game

Thurston Sexton, Max Yi Ren

1.1.2 Overview

- **Why human search strategy?**
- Bayesian Optimization (IBO)
- Inverse Bayesian Optimization (IBO)
- Case Study: Learning from others vs by one's-self
- IBO vs. Inverse Reinforcement Learning

1.1.3 Why human search strategy?

Human beings can **sometimes outperform** optimization algorithms at searching

- e.g. Figuring out how to play a game well.
 - Machine achieved *lower* than human level performance after 38 days of continuous play of ATARI Frostbite (Mnih et al.2015)
 - State-of-the-art deep learning achieves 3.5% of human performance on this game after the same amount of game play (2 hours)

Lake et al., “Building Machines That Learn and Think Like People”, *arXiv* (July 2016)
ecoracer.herokuapp.com

1.1.4 Why do humans succeed where algorithms fail?

- **Knowledge** - Prior information (or physical intuition)?
- **Meta Game** - ability in abstraction?
- **Learning to Learn** - Self improvement/strategic insight?

1.1.5 What about Crowdsourcing?

- **Currently:** wait for *good solutions* to arise by watching many players attempt solutions.
 - *a.k.a. smart people spending lots of time on a game*
- **What if:** wait for *good search strategies* by observing fast-improving players' decisions.
 - *a.k.a. smart people willing to at least try the game*

→ **Problem:** If we assume Humans search using an optimization scheme (e.g. Bayesian Optimization) with parameters informed by prior knowledge (their “strategy”), can we recover those parameters using their search trajectory?

1.1.6 Overview

- Why *human* search strategy?
- **Bayesian Optimization (BO)**
- Inverse Bayesian Optimization (IBO)
- Case Study: Learning from others vs by one's-self
- IBO vs. Inverse Reinforcement Learning

1.1.7 What is Bayesian Optimization?

Efficient Global Optimization (EGO) is excellent at optimizing expensive functions with minimal calls. 1. Sample the solution space 2. estimate/update the GP (kriging model) 3. Find new sample by maximizing a utility/aquisition function 4. repeat 2,3

Utility function can be probability of improvement, **expected improvement (EI)**, upper confidence bound, etc.

Search is therefore governed by GP parameters (λ) and the probabilistic nature of optimizing EI (strict vs non $\rightarrow \alpha$).

1.1.8 What is BO? Simple 1-D Example

$$\min_x f(x) \quad \text{where} \quad f(x) = e^{-(x-2)^2} + \frac{1}{10}e^{-(x-6)^2} + \frac{1}{x^2 + 1}$$

```
In [16]: # Our Objective Function
```

```
def target(x):  
    return np.exp(-(x - 2)**2) + np.exp(-(x - 6)**2/10) + 1/(x**2 + 1)  
  
x = np.linspace(-2, 10, 1000)  
y = target(x)  
bo = BayesianOptimization(target, {'x': (-2, 10)})  
  
# Gaussian Process Parameters  
gp_params = {"corr": "cubic", "theta0": 0.1, "thetaL": None, "thetaU": None}  
# Give the Model 3 random points to initialize  
bo.maximize(init_points=3, n_iter=0, acq='ei', xi=0.05, **gp_params)
```

Initialization

Step	Time	Value	x
1	00m00s	1.01744	5.6372
2	00m00s	1.31674	1.6754
3	00m00s	1.26393	2.4163

Bayesian Optimization

Step	Time	Value	x
------	------	-------	---


For the set of available samples \mathbf{X}_k and associated objective values \mathbf{f}_k ,

$$h_k := \langle \mathbf{X}_k, \mathbf{f}_k \rangle$$

Then given a utility function Q_{EI} ,

$$x_{k+1} = \operatorname{argmax}_{x \in \mathcal{X}} Q_{EI}(x; h_k, \lambda)$$

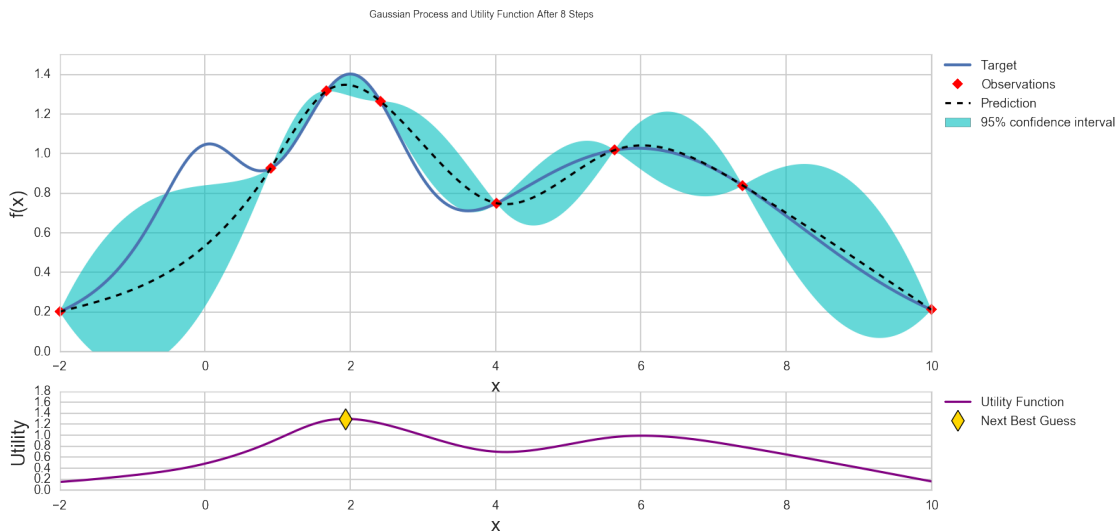
```
In [17]: plot_gp(bo, x, y)
```



```
In [22]: #Run until convergence
bo.maximize(init_points=0, n_iter=1, xi=.05) # random init, pts to add, p
plot_gp(bo, x, y)
```

Bayesian Optimization

Step	Time	Value	x
8	00m00s	0.83950	7.4019



1.1.9 What is BO? Pros/Cons

- fewer objective function calls
- creates “*surrogate model*” with cheap evaluation
- estimate of certainty is included
- matches actual human search in 1-D (*Borji & Itti 2013*)

but... - needs good Gaussian Process parameters for efficiency - curse of dimensionality
we'll be coming back to these

1.1.10 Overview

- Why *human* search strategy?
- Bayesian Optimization (BO)
- **Inverse Bayesian Optimization (IBO)**
- EcoRacer - Crowdsourced Strategy

We want to recover λ (and α) from an observed trajectory \mathbf{h}
Assume it's static. The Log-likelihood of a GP parameter can be defined as:

$$L(\lambda, \alpha) = \sum l(\alpha, \lambda, \mathbf{X}, \mathbf{f})$$

where l is the log-probability of every new sample *given* all previous samples/responses. **Every calculation of l needs integration over x**

1.1.11 IBO: Does it work?

Rosenbrock 30-D with four settings...how many samples do we need?

IBO quickly recognizes the true BO settings **unless** BO is behaving \sim random sampling (high λ)

1.1.12 IBO: Is it efficient?

Generally solvers find a λ at each iteration (adaptive) without prior knowledge (Maximum Likelihood Estimate)

1.1.13 Overview

- Why *human* search strategy?
- Bayesian Optimization
- Inverse Bayesian Optimization (IBO)
- **EcoRacer - Crowdsourced Strategy**
 - Dimension reduc
 - Results
 - IBO vs. Inverse Reinforcement Learning

1.1.14 EcoRacer

Part design (gear ratio), part control (trinary signal) Treated as a design optimization problem means **very high** dimensionality. - 18160 “ticks” where a decision can be made, spread over 36s - Nearly impossible to do BO without dimension reduction. - Previously: inputs like hill gradient, time since beginning, velocity, etc.

How do Humans “reduce the dimensionality” of a problem? *Salient features*, with minimal amount of multitasking.

→ **Independent Component Analysis (ICA)**

Unlike PCA (maximize variance in orthogonal bases), ICA performs the task of *blind source separation*. For an assumed number of *source signals* ICA finds sources such that they combine to produce the observed signals with **minimum shared information** (K-L divergence).

In []:

1.2 Numerical Modeling & Simulation

1.2.1 Solving PDE's

Ferromagnetic Phase Separation

MAE 598 (Spring 2016, Dr Yang Jiao)

CFD spinny thingies

MAE 561 (Fall 2015, Dr. Marcus Hermann)

In []: